

How Crowd Workers Use ChatGPT to Do Gig Work

Kiran Garimella
Rutgers University, USA

Abstract

Most evidence on how generative AI is reshaping work comes from outside the work itself: job postings that appear or vanish on freelance platforms, surveys that ask people whether they use AI, or controlled experiments. We instead observe the work directly, using a corpus earlier research has lacked: online crowd workers’ own ChatGPT conversation logs. Building on a dataset of 202,590 complete conversation exports donated by 1,252 Clickworker-recruited users in India, Nigeria, Brazil, and Pakistan over December 2022–February 2026 [16], we develop and validate a classifier for paid client (gig) work done with ChatGPT and analyze which tasks workers bring to the model, how much, how they collaborate with it, and how this changed. We report per-user estimands and relative measures rather than raw per-conversation counts. Four findings stand out. First, gig work is a real but bounded part of these workers’ ChatGPT use: it appears in about 10% of their conversations, and roughly four in five workers (79%) use the model for gig work at least once, though for the typical worker it accounts for about an eighth of conversations (mean 12%, median 8%); we report these as upper bounds and lead at the per-user level. Second, the paid tasks span the gig economy (writing, marketing, translation, design, coding, administrative work, tutoring), and the single largest category is data annotation, surveys, and microtasks, with crowdsourcing platforms named in the logs (Clickworker, Mechanical Turk, Appen); some workers use ChatGPT to complete the very microtasks and surveys that other platforms pay them for, which bears directly on the quality of crowdsourced data. Third, this share does not rise across three years; it holds steady even within individual users, while the task mix and the model generation turn over beneath that flat total. Fourth, gig use is highly concentrated across users (Gini 0.70) and is more common among ChatGPT Plus subscribers and younger workers. We discuss what this worker-side view implies for generative AI and precarious labor in the Global South.

1 Introduction

A large and growing share of paid work in the Global South is mediated by crowdsourcing platforms: short, piece-rate tasks such as data labeling, transcription, survey completion, content moderation, copywriting, and small coding jobs, distributed to a distributed workforce and paid by the task [1, 8]. Conversational AI arrived in the middle of this workday. The question this paper asks is simple to state and, until now, difficult to answer directly: when crowd and gig workers do their paid work, how do they actually use ChatGPT, and for which tasks, how much, and in what manner?

This is interesting and important for two reasons that pull in different directions. The first is a labor question. Crowd work is poorly paid (median wages on major microtask platforms have been estimated near \$2/hour, with much of the working day unpaid [1, 10]), and the four countries we study sit at the low end of that gradient. Sub-minimum piece rates create a strong incentive to finish faster, and a capable, cheap writing-and-coding assistant is exactly the tool that incentive selects for. Whether AI augments these workers (raising the throughput of a person who keeps the job) or automates them (doing work that would otherwise have employed another person) is a first-order question for hundreds of millions of people, and one that platform-side studies can only approach indirectly [5, 11]. The second is a data-quality question.

A large part of crowd work is producing the labels, transcriptions, and survey responses that researchers and companies buy as human data; when that work is done with a chatbot, the reliability of crowdsourced data is at stake [18, 20, 21]. We cannot observe what any individual task ultimately feeds, so we make the narrow version of this point rather than a broad claim about AI training data.

The problem is hard for three reasons. First, the behavior of interest is rare and ambiguously bounded: gig work is a thin slice of a corpus that is mostly personal use [16], and the line between doing a task *for a paying client* and doing one’s *own* income activity (a personal blog, an affiliate funnel, a small business) is genuinely subtle. Second, this slice is hard to detect automatically with precision, so estimates must be validated and read with care rather than taken at face value. Third, the natural quantities are statistically awkward: conversations are deeply unequal across users, so per-conversation averages are dominated by a few heavy contributors, and the data are retrospective histories truncated at each user’s export date, so a simple calendar trend confounds behavior change with changes in who is even present in the early months.

These difficulties explain why the question has not been answered directly. Platform-side studies observe job postings and prices, not what workers actually do [5, 11]; task-exposure indices score occupations

against capability checklists [7]; and laboratory and field experiments measure productivity on assigned tasks [2, 4, 12]. The line of work closest to ours infers LLM use inside crowd work statistically, by detecting machine-generated text in the responses workers return, without ever seeing the worker’s side of the interaction [18–21]. Conversation-level studies of AI use, including large platform reports [3, 9] and our own four-country dataset paper [16], describe usage broadly but do not isolate gig work. No prior study has had the crowd worker’s own gig conversations.

We do, and this paper uses them. We build on the corpus, demographics, and base classification of Roy Chowdhury and Garimella [16], and contribute: (i) a taxonomy and a two-stage classifier for gig work done with ChatGPT, distinguishing for-client work from a worker’s own job, study, and job search, and assigning a task type, a collaboration mode, and behavioral flags; (ii) a validation of those labels against a frontier model reading full conversation text, which we use to set how conservatively the numbers should be read; and (iii) an analysis establishing that the paid tasks span the gig economy, with microtask and survey work the single largest category and crowdsourcing platforms named in the logs; that the gig share is flat across three years even within users while the task mix and model generation change beneath it; and that gig use is concentrated and tied to subscription and age. We are explicit about what a single-platform convenience sample supports: the task mix and relational findings speak to the crowd-worker population from which the sample is drawn, while absolute prevalence and demographic composition describe this corpus only.

2 Related work

2.1 The economics of crowd and gig work

Online crowd workforces are young, comparatively well educated, geographically concentrated in the Global South, and economically precarious [1, 6, 14]. Hara et al. [10] estimate a median Mechanical Turk wage near \$2/hour and find that a large share of working time is unpaid, spent searching for and managing tasks. Berg et al. [1] document, across five platforms, that a substantial minority of workers rely on crowd work as a primary income and that pay varies sharply by region. Gray and Suri [8] situate this labor as the hidden human substrate of automated systems, what they call “ghost work”, and argue that its invisibility is structural rather than incidental. Our four countries occupy the low-pay end of this literature, which is the standing material incentive behind everything we observe.

2.2 Generative AI and the labor market

A fast-growing body of work measures how generative AI affects work. Task-exposure analyses estimate that a large fraction of occupations have significant shares of tasks exposed to LLMs, with knowledge-gig tasks such as writing, translation, and programming among the most exposed [7]. Field and laboratory experiments find large productivity gains that are concentrated among less-skilled workers and that compress the performance distribution [2, 12], alongside a “jagged frontier” where AI helps on some tasks and quietly hurts on others [4]. On the platform side, Hui et al. [11] find that demand and earnings for automation-exposed freelance work fell after ChatGPT’s release, with the highest-rated freelancers hit hardest, and Demirci et al. [5] report declines in automation-prone job posts. These studies are inferred from postings, prices, and controlled tasks; none observes the worker doing the work.

2.3 LLM use inside crowd work

The thread closest to our motivation studies whether crowd workers already use LLMs on the tasks they are paid for. Veselovsky et al. [18] combined a synthetic-text classifier with keystroke logging to estimate that 33–46% of Mechanical Turk workers used an LLM on a text-summarization task, and follow-on work studies the prevalence and prevention of such use [19]. Parallel studies detect AI-generated responses in crowdsourced surveys and on Prolific [20, 21], and industry and journalism report the same shift in freelance writing and academic ghostwriting [13, 15]. All of this evidence is recovered from the *output*: the text the worker returned, scored after the fact for machine authorship. We add the missing half of the interaction, the worker’s own prompts, which lets us see not only that a model was used but for what task, with what platform named, and in what manner.

2.4 Conversation-level studies of AI use

A final thread analyzes conversational AI use at the level of the conversation. Platform reports describe usage at enormous scale but release only aggregates over fixed taxonomies with inferred demographics [3, 9]. Research-proxy corpora such as WildChat make conversation-level data public but skew toward technically sophisticated users who self-select into a free research frontend [22], and privacy-preserving bottom-up systems summarize millions of conversations without releasing them [17]. Closest to us, Roy Chowdhury and Garimella [16] analyze the same four-country donated corpus we use, and report, among other findings, that personal use dominates (55–64% of conversations), that online-earning strategies form a recurring cluster

accounting for 10–17% of conversations in India and Pakistan, and that users have shifted over time toward delegating tasks to the model. That paper characterizes usage in general; it does not isolate gig work or validate a gig classifier. The present paper takes the earning-and-work fraction of that corpus as its subject.

3 Data

We use the corpus introduced and documented by Roy Chowdhury and Garimella [16], and refer the reader there for full detail on recruitment, consent, demographics, and the base classification pipeline; we summarize only what is needed here. The data are complete ChatGPT conversation exports donated by 1,252 users recruited and paid through Clickworker in India, Nigeria, Brazil, and Pakistan, totaling 202,590 conversations and spanning December 2022 to February 2026. Each conversation carries self-reported age and gender, a subscription tier, the model versions used, per-conversation length and language, and a set of classifications produced with frontier models following the OpenAI and Anthropic usage taxonomies, including a 24-category topic, an *Asking/Doing/Expressing* intent, and a *work/coursework/personal* task-purpose gate. The sample is a convenience sample of Clickworker-accessible, digitally literate ChatGPT users; following Roy Chowdhury and Garimella [16] we read absolute quantities as conditional on that frame and reserve generalization for relational findings and for claims about the crowd-worker population the sample is drawn from. The study is IRB-approved; third-party material in the logs (client names, pasted personal data, student identifiers) is removed before any text is read or quoted, and we report aggregates and paraphrased vignettes only. Our entry point into this corpus is the *work/coursework/personal* gate: we send the work-related, coursework-related, and ungated conversations (85,978 in total) to the gig classifier described next, and leave the personal majority aside, which makes our prevalence figures conditional on the work-related corpus and therefore lower bounds on gig work overall.

4 Methods

4.1 What counts as gig work, and a taxonomy

We define a gig-work conversation as one in which the user produces a deliverable that a third party (a client, an employer’s customer, or a crowdsourcing requester) will receive, submit, or pay for. The central and most error-prone distinction is between this and the user’s *own* income activity: running their own business, posting to their own monetized channel, or building an affiliate funnel are economic, but they are not for-client gig work, and the codebook places them under a separate *own job or business* label. We further sep-

arate the user’s *own study* (their coursework) and their *own job search* (CV, cover letter, interview preparation). Within the gig-for-client conversations we assign one of eleven task types, defined in Table 1 and grounded in the structure of crowd and freelance markets [1, 6], a collaboration mode adapted from the Anthropic Economic Index [9] that we collapse to automation-like use (delegating a task and taking the output) versus augmentation-like use (iterating, learning, or asking for validation), and four binary flags: whether the user writes in another person’s voice (impersonation), asks to evade an AI or plagiarism detector (detection-evasion), pastes a brief or rubric they did not author, and produces a usable artifact.

4.2 A two-stage classification funnel

Running a multi-field classifier over the full corpus is unnecessary because most conversations are personal. We therefore use a funnel. Stage 1 assigns the work-context label (gig-for-client, own job or business, own study, job application, or unclear) to the 85,978 work-related, coursework-related, and ungated conversations. Stage 2 assigns the task type, collaboration mode, platform, and flags, and is run only on the gig-for-client and unclear conversations, where the gig question is live. Both stages run locally with Gemma-3-27B (instruction-tuned), at temperature zero with JSON-constrained output, on a snippet of each conversation: the title and the first several user turns for Stage 1, and an extended snippet including an assistant turn for Stage 2, so the model can see the form of the deliverable. The funnel and the resulting label counts are given in Table 2; Gemma returned a valid Stage-1 label for 83,142 of the eligible conversations.

4.3 Validation against a frontier reference

Automatic detection of a rare, subtly bounded class is exactly the setting in which a classifier can look confident and be wrong, so we audit it. We drew a stratified sample of conversations (120 for Stage 1, 100 for Stage 2), oversampling the decision-critical gig-for-client and unclear classes and forcing coverage of all four countries and of the non-English (Portuguese, Hindi, Urdu, code-switched) rows. A frontier model (a current OpenAI model) then relabeled each sampled conversation under the identical codebook but reading the *full* conversation text rather than the snippet, which also lets us measure how much the snippet view misses. We treat the frontier labels as the reference and report agreement, per-class precision and recall, and the same for the Stage-2 fields (Table 3).

The audit is consequential. Against the reference, Gemma’s precision on the gig-for-client class is only 0.17, with recall 0.80: it finds most true gig work but labels far too much else as gig. Inspection of the false

positives shows they are overwhelmingly the user’s *own* income activity (affiliate marketing, a personal blog or YouTube channel, a small business), which the codebook assigns to own-job-business. The Stage-2 task type is moderately reliable (agreement 0.60) and the deliverable-produced flag is excellent (precision 0.98), but the collaboration mode is unreliable (agreement 0.48), and two flags are biased in known directions: detection-evasion is over-flagged (precision 0.33, recall 1.00) and impersonation is under-detected (recall 0.32). Agreement is similar across English and non-English rows (0.47 vs 0.47 overall), so the error is not a language artifact.

These results dictate how we report everything that follows. We treat Gemma as a high-recall filter rather than a precise labeler. Raw gig prevalence is reported as an upper bound, alongside a precision-scaled estimate (the raw rate times the measured positive predictive value); we do not apply a full sensitivity-specificity correction because the validation sample deliberately oversampled the gig classes, so its specificity is not representative of the corpus. Crucially, the task-mix results are *relative* shares within the flagged set and the flag directions are established, so they are far less sensitive to the precision problem than absolute counts are. A full frontier-model relabel of the candidate pool would convert the bounded prevalence into a point estimate and is the natural next step; the pool is already assembled for it.

4.4 Estimands and statistical approach

Conversations are nested within users who contribute very unequally, so we lead with two per-user quantities (the share of a user’s conversations that are gig work, averaged across users, and the share of users with any gig conversation), and report the pooled conversation-level share alongside. All regressions are estimated at the user level. We summarize concentration with the Gini coefficient and Lorenz curve of gig conversations across users. For the time trend we use a within-user (fixed-effects) transformation: we subtract each user’s mean before estimating the monthly trend, so the result reflects change inside users rather than change in the composition of who is observed each month. This matters because the exports are retrospective histories truncated at each person’s recruitment date, so the earliest months contain only early adopters. We model whether a user does any gig work with logistic regression on age, gender, country, and subscription tier, and gig intensity with a negative-binomial model using the log of the user’s conversation count as an exposure offset. We trim months with fewer than 20 conversations and report odds and incidence-rate ratios with confidence intervals.

Table 1: Taxonomy of gig task types assigned in Stage 2.

Task type	Definition
data annotation	labeling, tagging, transcription, filling surveys, microtasks, HITS
writing content	articles, blog posts, SEO copy, product descriptions
marketing/social	social posts, ad copy, campaigns, SEO strategy
design/media	image, graphic, video, or audio production
academic ghost	essays or assignments written for a student/client
translation	translating or localizing text
coding/dev	software, scripts, web, debugging for a client
admin/VA	data entry, scheduling, email, customer service
tutoring others	teaching materials or answers produced for others
proposal/bid	freelance proposals, bids, cover letters, client messages
other gig	for-client paid work not covered above

Table 2: Classification funnel. Stage 1 is run on the work-related, coursework-related, and ungated conversations; Stage 2 on the gig-for-client and unclear conversations.

Stage / label	Conversations
Corpus [16]	202,590
personal (set aside)	116,552
Stage-1 eligible (work/coursework/ungated)	85,978
Stage-1 labeled by Gemma	83,142
gig-for-client	20,827
own job or business	11,649
own study	26,001
job application	4,577
unclear	20,088

5 Results

5.1 How much gig work, and how concentrated

On the raw Gemma labels, 10.3% of conversations are gig-for-client and 78.9% of users have at least one such conversation; the mean per-user gig share is 12.2% (median 7.7%). These raw classifier rates are upper bounds (Methods); a precision-scaled version places the conversation-level figure near 2% (Table 4). The takeaway is not the exact rate but the shape: gig work touches many workers occasionally yet constitutes a small and concentrated part of what any of them does. The Gini of gig conversations across users is 0.70, and the top decile of users accounts for 52% of all gig conversations (Figure 3). This concentration is why we avoid a per-conversation headline; a handful of heavy users would otherwise set it.

Table 3: Validation against a frontier model reading full conversation text. Stage 1 $n=120$, Stage 2 $n=100$; gig and rare classes oversampled.

Stage	Quantity	Value
1	overall agreement	0.47
1	gig-for-client precision	0.17
1	gig-for-client recall	0.80
1	agreement, English / non-English	0.47 / 0.47
2	task type agreement	0.60
2	collaboration mode agreement	0.48
2	deliverable-produced precision	0.98
2	detection-evasion precision / recall	0.33 / 1.00
2	impersonation precision / recall	0.60 / 0.32

5.2 What kind of gig work

Among gig conversations, the task mix (Figure 1, left; by country in Figure 5) is led by data annotation and microtasks (20.3%), content writing (15.2%), and marketing or social copy (14.1%), followed by design and media (9.1%), admin/virtual-assistant work (7.7%), tutoring for others (7.5%), academic ghostwriting (7.4%), and coding (6.1%); translation and proposal writing make up the tail. The ordering is stable across the four countries, with modest shifts (for example, more coding in India, more translation in Pakistan), consistent with the country-level specializations reported by Roy Chowdhury and Garimella [16].

5.3 Microtask and survey work

The single most common gig task is data annotation, surveys, and microtasks, and the crowdsourcing platforms that distribute this work appear by name inside the conversations: Clickworker (873 conversations), Mechanical Turk (500), and Appen (296) (Figure 1, right). That is, some of these workers use ChatGPT to complete the very microtasks and surveys that other platforms pay them for. We cannot tell what any individual task ultimately feeds, so we do not characterize it beyond this; the narrow and well-supported implication is for the quality of crowdsourced data. The LLM use that Veselovsky et al. [18] and Zhang et al. [21] recover statistically from returned text is here visible from the worker’s own side, with the platform named and the task in view, direct evidence that crowdsourced labels and survey responses are, in some share, produced with a chatbot.

How does the task reach the model? Most often the worker pastes it in. In 28% of gig conversations the user pastes a brief, rubric, job post, or task they did not author (Figure 6), and this rises to 44–45% for coding, academic ghostwriting, and data-annotation work, where a discrete external specification is the norm. The flag is over-detected in validation, so 28% is an upper

bound, but the gradient across task types is clear and matches what those tasks look like.

Two related behaviors appear at lower measured rates, and the validation tells us which way each is biased. Among gig conversations, 3.5% involve impersonation (answering in an assigned persona or another person’s voice) and 1.4% ask the model to evade an AI or plagiarism detector. Because the classifier under-detects impersonation (recall 0.32), 3.5% is a floor; because it over-flags detection-evasion (precision 0.33), 1.4% is a ceiling. The qualitative cases are unambiguous: persona-conditioned survey answering, and requests to rewrite text so that it passes a detector.

5.4 How workers collaborate with the model

Collapsing the collaboration mode, gig conversations split almost evenly between automation-like use (51%) and augmentation-like use (49%), shown by task type in Figure 7. We report this descriptively only: validation agreement on this label is 0.48, so the split should be read as indicative. It is worth stating because Roy Chowdhury and Garimella [16] find a corpus-wide drift toward task delegation over time; within gig work specifically, both modes are well represented, and the delegation share does not crowd out iterative and learning-oriented use.

5.5 Has it changed over three years?

It has not, in the way one would expect. The gig share is flat (Figure 8, top): it holds near 10% of conversations from 2023 to 2026, and the within-user fixed-effects slope is effectively zero (0.0006 per month). The raw calendar series and the within-user series track each other closely, so the flatness is not an artifact of changing composition. What moves sits beneath the total. The task mix shifts steadily (Figure 8, bottom), and the shift lines up with the model generation in use (Figure 9): data annotation roughly doubles from the GPT-3.5 to the GPT-5 generation (11% to 25% of gig conversations) while content writing falls (21% to 11%) and marketing and academic ghostwriting decline as well, a difference that is highly significant (a χ^2 test comparing the GPT-4o and GPT-5 generations gives $p < 10^{-38}$). Collaboration tilts the same way, from automation toward augmentation as the models improve (60% to 46% automation-like across the three generations), though that label is our least reliable. Because model generation is collinear with calendar time, we read this as the changing composition of a stable total rather than a causal effect of capability: the same amount of gig work, done with more capable models, on a shifting set of tasks.

5.6 Who does gig work

A user-level logistic model of any gig use (Figure 4) shows that ChatGPT Plus subscribers are about $2.5\times$ as likely to do gig work (odds ratio 2.54, $p=0.001$) and that the likelihood declines with age (odds ratio 0.975 per year, $p=0.002$). Relative to Brazil, the odds are significantly lower in India (0.38, $p<0.001$) and Pakistan (0.50, $p=0.007$) and indistinguishable in Nigeria. Gender is not significant in the model, although women show a slightly higher share descriptively. The negative-binomial intensity model points the same way, with Plus status and younger age associated with more gig conversations per user. The Plus association is notable: in this corpus the \$20/month subscription represents a meaningful fraction of monthly income [16], so the workers who pay for the more capable model are also the ones who lean on it for paid work, which fits a capability-and-investment reading rather than a pure-precarity one.

5.7 The scale of output

As an order-of-magnitude indicator, gig conversations produced a large volume of model text, with a median near 1,340 words of assistant output per conversation; the distribution by task type is in Figure 10. Attaching an illustrative rate of \$5–20 per thousand words to the text-producing tasks gives a sense of the dollar scale of the assisted output, but we stress this is a scale indicator and not an earnings estimate: we observe no payments, and much crowd work is paid per task or per hour rather than per word.

6 Discussion

A flat share is not a flat amount. The result most easily misread is the flat trend, so it is worth being exact about what stays constant. Within a given worker, the *fraction* of conversations that is gig work does not grow from 2023 to 2026. The worker’s total use of ChatGPT does grow over that window as it spreads into personal and everyday tasks. A constant fraction of a growing whole means the absolute amount of gig work done with the model is rising even though its share is not. What is stable is the place gig work holds in a worker’s relationship with the tool, not the volume of it. That reframes what the rest of this section is about: under a roughly fixed gig share, the interesting action is in the composition underneath it and in who is doing the work.

The composition shift echoes the demand shift seen from the platform side. The clearest movement in the data is compositional. As a share of gig conversations, content writing, marketing copy, and academic ghostwriting fall while data annotation and microtasks rise (Figure 9). That is the same direction platform studies report from the outside: demand and earnings for

automation-exposed freelance writing dropped after ChatGPT, and the most exposed work was hit hardest [5, 11]. The two vantage points line up into one account. Clients who once hired a person to write now prompt a model themselves, so the writing gigs thin out, and the work that remains worth posting is the labeling, transcription, and survey work that still needs a person attached to it. We cannot fully separate this demand-side reading from a simpler supply-side one, in which annotation is merely more amenable to a chatbot than it used to be, because model generation moves with calendar time. Even so, it is notable that what a worker brings to ChatGPT shifts in step with what the market outside is doing to their options.

Pasting the task and taking the output is the mechanism, not a detail. Three findings combine into a single account of how crowdsourced data comes to be model-generated. Annotation and survey work is the largest gig category, it leans toward delegation rather than iteration (Figure 7), and the task is most often pasted in verbatim, with copy-paste highest precisely for annotation, coding, and academic work (Figure 6). Put those together and the route is concrete: a worker pastes the requester’s item, asks for the answer, and submits it, so the human judgment the requester pays for is in that moment the model’s reply to the requester’s own prompt. This is why our data-quality caution is narrow but firm. The risk is not spread evenly across crowd work. It concentrates where a discrete external specification meets a delegation-style workflow, and that combination is measurable in advance, so a buyer of crowdsourced labels can tell which task types are most exposed and put gold-standard and detection checks there first. We make no claim about what any task feeds, only that anyone relying on crowdsourced human responses should expect a share of them to be model-assisted.

Who captures the gains points to stratification, not only compression. Experiments have stressed that generative AI narrows skill gaps, helping less experienced workers most [2, 12]. Our cross-section adds a countervailing pattern. Gig use is concentrated, with a tenth of workers producing half of it, and the strongest correlate of using ChatGPT for paid work is paying for it (Figure 4). Where a \$20 monthly subscription is a genuine expense, the workers who can invest in the more capable model are also the ones who turn it into income. If the tool’s advantage is largest at the frontier model, then access to that model, gated by ability to pay, can reintroduce the inequality that compression is supposed to remove. We cannot test this without earnings, but the joint pattern of concentration and the subscription effect is exactly what a distributional account of AI and gig labor should be looking at, and it cuts against an unqualified optimism.

What the worker’s-eye view does and does not set-

tle. On the augmentation-versus-automation question, the honest reading of these logs is that both occur, often in the same person, who delegates some tasks whole-sale and iterates, learns, or checks their own work on others. The balance tilts toward iteration as the models improve, though this is our least reliable measure and we hold it loosely. For a workforce at the bottom of the global pay scale, the hopeful case, in which a worker keeps the job and produces more, and the harsher case, in which the model does work that would have employed someone else, are not abstractions to be settled by argument; they coexist in the data, and their balance is an empirical question. What conversation logs offer that postings and surveys cannot is the act itself: the task, the platform, the prompt, and the way the model is asked to help. What they withhold is the outcome, the pay, the rejection, the next gig. Linking this worker-side view to earnings and to platform-side demand is the natural next step, and the donated-log method makes it feasible.

7 Limitations

The sample is a convenience sample of Clickworker crowd workers in four countries; absolute prevalence and demographic composition describe this corpus and should not be read as country-level estimates, which is why we lead with per-user shares and relative measures. As detailed in Methods, automatic detection of for-client work is imperfect, so we treat raw rates as upper bounds and report a precision-scaled prevalence alongside; a full frontier-model relabel of the candidate pool, which we have prepared, would convert these into point estimates. The collaboration-mode label is the least reliable and we treat it as suggestive. Snippet-based classification under-detects gig work whose for-client context appears late, so prevalence is a lower bound conditional on the work-related conversations sent to the classifier, and gig work among conversations the upstream gate labeled personal is not counted at all. The economic figures are illustrative; we observe no payments. Finally, the validation is agreement with a frontier model, not human ground truth; a human-labeled gold set would strengthen it and is straightforward to add.

8 Conclusion

With crowd workers’ own ChatGPT logs we can see what platform-side and laboratory studies could only infer. Workers in India, Nigeria, Brazil, and Pakistan bring a wide range of paid client tasks to the model, from writing and translation to coding and, most commonly, microtask and survey work. The share of their conversations devoted to gig work has not grown over three years; the models they use have. Seeing this from the worker’s own side, rather than from postings or

Gig prevalence by country (raw classifier labels; read as upper bounds)




Figure 2: Gig-work prevalence by country. Per-worker mean gig share (bars, with 95% confidence intervals) and the conversation-level share overlaid as points, plus the share of workers with at least one gig conversation. These are raw classifier labels, read as upper bounds (Methods); we lead with the per-worker quantities.




Figure 3: Gig work is concentrated in a small set of workers. The Lorenz curve of gig conversations across the 1,252 workers shows that the most active tenth produce 52% of all gig conversations (Gini 0.70). For most workers gig work is occasional, and for a few it is routine, which is why we report per-worker rather than per-conversation quantities.

returned text, is what lets us state it plainly, and what makes the data-quality and labor implications concrete.

References

- [1] J. Berg, M. Furrer, E. Harmon, U. Rani, and M. S. Silberman. Digital labour platforms and the future of work: Towards decent work in the online world. Technical report, International Labour Office, 2018.
- [2] E. Brynjolfsson, D. Li, and L. R. Raymond. Generative AI at work. *Quarterly Journal of Economics*, 2025.
- [3] A. Chatterji, K. Handa, et al. How people use ChatGPT. Technical report, National Bureau of Economic Research, 2025.




Figure 1: What gig work workers bring to ChatGPT. Left: the distribution of gig task types among conversations flagged as for-client work, ordered by frequency, with data annotation (red) the single largest at 20%. Right: the number of gig conversations that name a crowdsourcing platform. Shares are relative within the flagged set, which makes them robust to the classifier’s precision (Methods); the ordering and the named platforms, rather than the absolute level, are the load-bearing facts.




Figure 4: Who uses ChatGPT for gig work. Odds ratios from a worker-level logistic regression of any gig use on age, gender, country, and subscription, with 95% confidence intervals on a log scale (* $p < .05$, ** $p < .01$, *** $p < .001$). Paying for ChatGPT Plus is the strongest correlate at $2.5\times$, older workers are slightly less likely, and India and Pakistan sit below Brazil. Country effects describe this Clickworker sample, not the national populations.




Figure 5: Gig task mix by country (composition within each country’s gig conversations; same task colors and order as Figure 1). The leading categories are stable across countries, with smaller country-specific tilts such as more coding in India and more translation in Pakistan.

- [4] F. Dell’Acqua, E. McFowland, E. Mollick, H. Lifshitz-Assaf, K. Kellogg, S. Rajendran, L. Kraye, F. Candelon, and K. R. Lakhani. Navigating the jagged technological frontier: Field experimental evidence of the effects of AI on knowledge worker productivity and quality. Technical report, Harvard Business School Working Paper 24-013, 2023.
- [5] O. Demirci, J. Hannane, and X. Zhu. Who is AI replacing? the impact of generative AI on online freelancing platforms. *Management Science*, 2025.
- [6] D. Difallah, E. Filatova, and P. Ipeirotis. Demographics and dynamics of Mechanical Turk workers. In *Proceedings of the Eleventh ACM International Conference on Web Search and Data Mining (WSDM)*, 2018.
- [7] T. Eloundou, S. Manning, P. Mishkin, and D. Rock. GPTs are GPTs: Labor market impact potential of large language models. *Science*, 2024.
- [8] M. L. Gray and S. Suri. *Ghost Work: How to Stop Silicon Valley from Building a New Global Underclass*. Houghton Mifflin Harcourt, 2019.
- [9] K. Handa et al. The Anthropic economic index. Technical report, Anthropic, 2025.
- [10] K. Hara, A. Adams, K. Milland, S. Savage, C. Callison-Burch, and J. P. Biggam. A data-driven




Figure 9: What workers do, and how, shifts across model generations. Left: among gig conversations, the share of the six task types that change most, by the model generation in use; data annotation roughly doubles from the GPT-3.5 to the GPT-5 generation while writing, marketing, and academic ghostwriting fall. Right: the automation versus augmentation split across the same generations. Model generation is collinear with calendar time, so these are descriptive shifts rather than a causal effect of capability, and the collaboration split is the least reliable label (Methods).




Figure 6: Workers often paste the client's task into ChatGPT. The share of gig conversations in which the user pastes a brief, rubric, job post, or task they did not author, by task type (overall 28%, dashed line; bars at or above 40% in red). The behavior is most common for coding, academic ghostwriting, and data-annotation work, where a discrete external specification is the norm. Read as an upper bound (Methods).

Figure 7: Collaboration mode by gig task type. Each task's gig conversations are split into automation-like use (delegate the task and take the output) and augmentation-like use (iterate, learn, or validate); tasks are in the same order as Figure 1. Translation and data annotation lean most toward delegation. This collaboration label is our least reliable (validation agreement 0.48, Methods), so the figure is indicative rather than precise.

analysis of workers' earnings on Amazon Mechanical Turk. In *Proceedings of the 2018 CHI Conference on Human Factors in Computing Systems*, 2018.

- [11] X. Hui, O. Reshef, and L. Zhou. The short-term effects of generative artificial intelligence on employment: Evidence from an online labor market. *Organization Science*, 2024.
- [12] S. Noy and W. Zhang. Experimental evidence on the productivity effects of generative artificial intelligence. *Science*, 2023.
- [13] Originality.AI. Analysis of AI-generated content among Upwork writers, 2025.

- [14] L. Posch, A. Bleier, F. Flöck, and M. Strohmaier. Characterizing the global crowd workforce: A cross-country comparison of crowdworker demographics. *Human Computation*, 2018.
- [15] Rest of World. Kenyan academic ghostwriters and the adoption of ChatGPT, 2023.
- [16] S. Roy Chowdhury and K. Garimella. How people use ChatGPT: Conversation-level evidence from India, Nigeria, Brazil and Pakistan. *Manuscript; data and code at <https://github.com/gurkiran/how-people-use-chatgpt-data>*, 2026.




Figure 8: The gig-work share stays flat while its composition moves. Top: the monthly share of conversations that are gig work, shown as the raw calendar mean (grey) and after subtracting each worker’s own mean so the line reflects within-worker change (red); both are flat across 2023 to 2026, with two model-release dates marked. Bottom: the quarterly composition of gig task types (same colors and order as Figure 1).

Table 4: Gig prevalence by country: raw conversation share, precision-scaled estimate, per-user mean share, and share of users with any gig conversation.

Country	Conv. (raw)	Conv. (scaled)	Per-user mean	Any-gig users
Brazil	0.120	0.020	0.180	0.858
India	0.087	0.015	0.095	0.740
Nigeria	0.115	0.019	0.127	0.852
Pakistan	0.108	0.018	0.121	0.767




Figure 10: Illustrative output volume by gig task type, measured as the total assistant words produced. The dollar band applies a \$5–20 per-thousand-word rate to the text-producing tasks as a scale indicator only; we observe no payments, and much crowd work is paid per task or per hour.

[22] W. Zhao, X. Ren, J. Hessel, C. Cardie, Y. Choi, and Y. Deng. WildChat: 1M ChatGPT interaction logs in the wild. In *International Conference on Learning Representations (ICLR)*, 2024.

- [17] A. Tamkin, M. McCain, K. Handa, et al. Clio: Privacy-preserving insights into real-world AI use. *arXiv preprint arXiv:2412.13678*, 2024.
- [18] V. Veselovsky, M. H. Ribeiro, and R. West. Artificial artificial artificial intelligence: Crowd workers widely use large language models for text production tasks. *arXiv preprint arXiv:2306.07899*, 2023.
- [19] V. Veselovsky, M. H. Ribeiro, P. Cozzolino, A. Gordon, D. Rothschild, and R. West. Prevalence and prevention of large language model use in crowd work. *Communications of the ACM*, 2025.
- [20] S. Xu, Y. Zhang, and A. J. Alvero. Reliable or deceptive? investigating AI-generated survey responses on Prolific. Working paper, 2024.
- [21] Y. Zhang, T. Katoh, and J. Pei. Detecting the use of generative AI in crowdsourced surveys. Working paper, 2024.