

The Hype vs. the Reality of Generative AI in Elections: The Case of the 2024 Indian Elections

Kiran Garimella*

Simon Chauchard[†]

Accepted for publication in *The International Journal of Press/Politics*

This is the authors' accepted manuscript. The version of record will appear at the journal.

Abstract

The 2024 Indian General Election was defined by pervasive warnings that generative AI would unleash a tsunami of deepfakes and fundamentally destabilize democratic processes. To empirically test these predictions, we analyzed over 5.5 million messages from more than 6,000 WhatsApp groups, obtained through a large-scale data donation project in India. Our findings reveal a profound disconnect between pre-election hype and ground truth: generative AI constituted only about 1.1% of the image dataset once corrected for the measured recall of our detection pipeline, a figure we treat as a lower bound, with no evidence of a viral “deepfake” epidemic. Furthermore, a qualitative taxonomy of this content reveals that the technology was rarely deployed for deception. Instead, it was primarily utilized for “phatic” labor—automating the production of hyper-stylized “Good Morning” greetings, religious iconography, and patriotic art. We argue that AI adoption was constrained by technological friction and by economic incentives, as traditional “cheapfakes” remained comparatively more cost-effective alternatives for political actors. These results challenge the prevailing focus on the threat of AI-driven messaging, suggesting that the technology may not *always* lead to epistemic disruption, but rather to the efficient reproduction of existing social rituals.

Keywords: Generative AI, WhatsApp, India, Elections, Propaganda.

*Library and Information Science, Rutgers University, United States. kiran.garimella@rutgers.edu. Both authors contributed equally to this work.

[†]Political Science, Universidad Carlos III de Madrid, Spain. simon.chauchard@uc3m.es

1 Introduction

The 2024 global election cycle was widely heralded as the “year of generative AI,” a turning point where synthetic media would fundamentally alter the mechanics of democracy. In the months preceding India’s general election—the world’s largest democratic exercise with over 900 million eligible voters—academic experts, security agencies, and technology vendors coalesced around a narrative of “enchanted determinism” (Campolo and Crawford 2020). Predictions ranged from a tsunami of undetectable deepfakes to the automated micro-targeting of propaganda at a scale beyond human comprehension (Cybersecurity and Agency 2024; Elliott 2024). This discourse posited that generative AI would act as an unstoppable, exogenous shock, rendering traditional safeguards of electoral integrity obsolete (Bueno de Mesquita et al. 2023). Accompanying this anxiety was a surge in commercial activity, with political consultancies and startups marketing AI-powered solutions for “personalized voter outreach,” creating a feedback loop that reinforced the perception of AIs inevitable omnipresence (Rebelo 2024). Even if empirical work on visual misinformation has repeatedly shown that threat perceptions outpace observed prevalence (Brennen, Simon and Nielsen 2021), citizens shared this anxiety: recent survey evidence from the 2024 European Parliament elections shows that majorities of respondents view generative AI as a threat to democratic processes, especially with regard to visual disinformation (Farooq et al. 2025).

However, despite the highly anxious tone of these predictions, few datasets exist to assess the empirical reality of the 2024 election. While *theoretical* models of AI harm abound, ecological data from active political environments is scarce. Understanding the actual influence of generative AI is uniquely difficult in the Global South, where political communication occurs primarily on end-to-end encrypted platforms like WhatsApp. Unlike public networks (e.g., X or Facebook) where API access allows for broad monitoring, WhatsApps dark architecture obscures the provenance and spread of content. Consequently, the existing literature suffers from a “streetlight effect”, focusing on few remarkable examples drawn from public platforms or small-scale qualitative interviews, leaving a profound gap between our anxieties about AI and our knowledge of its ground truth.

This paper addresses that gap by providing the first large-scale, privacy-preserving em-

pirical assessment of generative AIs actual prevalence in the Indian electoral ecosystem. We leverage a unique dataset of over 5.5 million messages collected from 6,000+ WhatsApp groups via a data donation infrastructure deployed during the 2024 general elections campaign across the Northern Indian state of Uttar Pradesh, the largest Indian state and a key political battleground in these elections. To overcome the challenges of detection at scale, we developed a rigorous multi-stage pipeline combining open-weights models (Gemma), state-of-the-art vision language models (GPT-4o), and expert human annotation. This approach allows us to move beyond simple keyword searches or unverified user reports, distinguishing high-fidelity deepfakes from the benign digital art that characterizes Indian digital culture. We further validate the recall of this pipeline’s open-weights first-pass filter, using both an external benchmark and an in-corpus audit, to bound how much AI content it might miss (Appendix D).

Our findings reveal a stark disconnect between the pre-election hype and the empirical reality. Contrary to predictions of an AI takeover or tsunami, generative AI content constituted only about 1.1% of the total image dataset, a recall-corrected estimate that we treat as a lower bound. Even among highly viral content, where one would expect effective propaganda to cluster, AI-generated media remained a negligible minority compared to traditional, human-generated political messaging. Furthermore, the nature of this content qualitatively defied the “disinformation” frame. Rather than malicious deepfakes designed to deceive, the vast majority of AI images were “phatic” in nature: hyper-stylized “Good Morning” greetings, religious iconography, and patriotic digital art that may have led to connective action (Bennett and Segerberg 2013), but did not neatly map into a partisan campaign strategy.

In addition to providing us with a reality check about the volume and the share of content that was AI-generated, these results also reveal what the main driver for the creation of AI content was during this electoral period. Namely, they suggest that the primary driver of AI usage during the 2024 electoral cycle was not the desire to generate more manipulative political content, but rather *efficiency*. Political actors and ordinary users alike appropriated the technology to automate the production of already common, low-stakes social content. In that sense, they fitted AI into the existing aesthetic vernacular of Indian WhatsApp groups instead of using it to create fundamentally different and disruptive types of content. We further identify significant technological friction and economic incentives as limiting factors: for political actors in charge of

creating political propaganda (“IT cells”), traditional “cheapfakes” and human-generated rumors remained more cost-effective and/or persuasive than the computationally expensive and widely distrusted generative AI productions available at the time.

While generative AI technologies have advanced rapidly since mid-2024, these findings nonetheless provide us with a crucial data point to understand how AI-generated images and videos may be used in future elections. The limited impact observed here may merely be a function of the technical immaturity of models like DALL-E 3 or Midjourney v6, which users would have relied on at the time. However, this study establishes a vital sociotechnical baseline: it demonstrates that the constraints on AI disruption are not merely technical, but *sociological* and *economic*. The failure of AI to achieve viral dominance in 2024 highlights the resilience of human distribution networks and the “authenticity paradox,” where the glossy aesthetic of AI triggers user skepticism. As models improve towards photorealism, these human dynamicstrust, distribution, and economic utilitywill continue to serve as the primary bottlenecks to adoption. Understanding why the “dog didn’t bark” in India in 2024 provides the necessary blueprint for distinguishing between genuine future threats and the persistent noise of enchanted determinism. Subsequent campaign-tracking work in the 2025 German and Dutch electoral environments documents that the volume and visual quality of AI-generated political content has since grown (Votta and Kruschinski 2025). Our study does not contest this trajectory; rather, it provides a baseline against which those more recent upticks can be interpreted, and it specifies the sociological, economic, and platform-level mechanisms that mediate whether improved technical capacity translates into widespread use during electoral campaigns.

2 How Might AI Affect Elections? Theoretical Expectations

The discourse surrounding generative AI in elections has been defined by a sharp tension between theoretical anxiety and empirical ambiguity. This section traces the trajectory of this debate, moving from the “supply-side” predictions of epistemic collapse to the emerging “demand-side” realities of technological friction, economic incentives, and cultural practice in the Global South.

The 2024 global election cycle, dubbed the “super election year”, was widely framed

through the lens of what Campolo and Crawford term *enchanted determinism*: a discourse where AI is portrayed as an autonomous, unstoppable force that operates beyond human agency (Campolo and Crawford 2020). Security and policy literature heavily amplified this narrative. The Munich Security Conference characterized the period as “AI-pocalypse Now?”, predicting that the marginal cost of creating high-fidelity disinformation would approach zero, leading to a tsunami of synthetic propaganda (Carr and Köhler 2024). Scholars argued that the modal richness of AI and its ability to generate photorealistic video and audio, would bypass traditional skepticism, creating a post-truth environment where shared reality is fundamentally fractured (Christopher and Bansal 2024). This “supply shock” thesis assumed that because the *capacity* to create deepfakes existed, the *deployment* would be inevitable and ubiquitous (Ferrara 2024).

Countering the determinist narrative, a growing body of literature on the political economy of disinformation suggests that propagators are rational economic actors constrained by technological friction. Paris and Donovan introduced the concept of “cheapfakes”—manipulations achieved through accessible tools like splicing, context-shifting, or speeding up video—arguing that these low-tech methods often yield higher returns on investment than computationally expensive deepfakes (Paris and Donovan 2019). Empirical work supports this cheapfake preference. Vaccari and Chadwick (2020) and Dobber et al. (2021) found that deepfakes do not necessarily alter voter attitudes more effectively than textual misinformation or simple edits. Furthermore, qualitative research indicates that audiences often view the ‘synthetic perfection’ of AI-generated content with suspicion, preferring ‘analog authenticity’ where imperfections serve as markers of truth. Participants expressed concern that unlike traditional photography, the limitless potential of AI to fabricate visuals threatens the concept of an objective, shared reality (Strikovic and Cools 2025). This literature provides a theoretical basis for our finding that political actors in India continued to prioritize human-generated content: the *cost* of upgrading to AI outweighed marginal potential benefits in terms of *persuasion*.

While the direct impact of deepfakes has been limited, scholars identify a secondary, perhaps more pernicious effect: the Liars Dividend. Schiff, Schiff and Bueno (2025) demonstrate that the mere *public awareness* of generative AI enables political actors to dismiss authentic scandals as fake news or deepfakes. This creates an ‘authenticity paradox’ where the threat of AI causes more epistemic damage than the AI itself. However, this dividend has a limit. Schiff et al.

find that this strategy is highly effective against text but faces resistance against video evidence, suggesting that the public retains a residual trust in visual evidence that 2024-era AI has not yet fully eroded. This dynamic may help explain the low prevalence or the low virality of AI content in social media datasets; as awareness of the technology spreads, the distinctive glossy aesthetic of Midjourney or DALL-E images may now function as a heuristic warning signal—a digital stigma—rather than a tool of deception.

Finally, to understand the prevalence of phatic content such as benign “Good Morning” AI images —common in the empirical context we focus on in this article, India—, we must look to the anthropology of digital communication in the Global South. Scholars of the Global South have long noted that platforms like WhatsApp are not just information conduits but infrastructures for “phatic labor”—the social work of maintaining relationships through small, non-informational gestures (Miller et al. 2016). In the Indian context, the “Good Morning” message is a massive cultural phenomenon, representing a form of digital care that bridges the gap between older, non-English speaking demographics and the internet (Escoe, Martin and Salerno 2025). While elite discourse often dismisses this content as spam, anthropological accounts frame it as a vital social glue. Our findings align with Daepp et al. (2025), who observe that creators in India often use AI not for deception (deepfakes) but for spectacle and efficiency using tools to quickly generate the hyper-saturated, devotional, or patriotic imagery that fuels this phatic economy. Besides, beyond mere citizens, political actors themselves appear to frequently post such content, likely in an effort to cultivate their audience (Chauchard and Garimella 2022). Following this logic, AI need not be weaponized to attack opponents, and is more likely to be domesticated to serve social media users’ relentless demand for social connection.

While existing literature has successfully mapped the *potential* for AI disruption and the *mechanisms* of its theoretical efficacy, there remains a critical paucity of ecological data from active political environments. Most current studies rely on small-scale qualitative interviews (Daepp et al. 2025), platform-specific data restricted to the Global North (Chen et al. 2025), or theoretical modeling of threat vectors (Ferrara 2024). Our work bridges this gap by providing the first large-scale, privacy-preserving empirical assessment of generative AIs’ actual prevalence in the world’s largest democracy. By shifting the unit of analysis from *what could happen* (theoretical capability) to *what actually happened* (empirical prevalence), this study offers a necessary corrective to the

enchanted deterministic narratives that have dominated the field, re-grounding the discourse in the material realities of friction, cost, and cultural practice.

2.1 Scope conditions and platform affordances

Because platform affordances, media economies, and political-communication repertoires vary across contexts, the patterns we document are bounded by a set of scope conditions that merit explicit articulation. First, diffusion of a new communicative technology is a social process: adoption depends on local network thresholds and on whether users observe trusted alters using the tool (Valente 1996). In an information ecosystem that rewards the personalized, text-heavy circulation of memetic content (Bennett and Segerberg 2013), the marginal return on a glossy synthetic image is low when non-AI strategies (e.g. “cheapfakes”) already satisfy users’ communicative needs (Freelon and Wells 2020; Guess, Nyhan and Reifler 2020). Second, platforms “digital architectures” structure what campaigns can and will produce (Gillespie 2010; Bossetta 2018). WhatsApp’s peer-to-peer, group-mediated, end-to-end-encrypted design rewards forwardable, low-friction content created or curated by trusted nodes rather than broadcast-style synthetic spectacle. Third, political communication in contemporary democracies operates through a hybrid media system in which legacy, digital, and interpersonal logics co-constitute one another (Chadwick 2017); in India, WhatsApp sits inside a hybrid ecology that still privileges television, vernacular press, and large-scale rallies as primary agenda-setters. Our findings on the limited prevalence and non-deceptive uses of AI therefore speak most directly to: (i) closed, group-centric messenger platforms; (ii) political systems in which a dominant party has already invested in large-scale human distribution networks; and (iii) periods during which generative models still produce visibly “glossy” artifacts that function as skepticism heuristics. The extent to which these findings generalize to broadcast channels, to less-organized partisan environments, or to a next generation of near-photorealistic models is an empirical question we return to in the limitations section.

3 Methods

3.1 Dataset

To empirically assess the prevalence of generative AI content, we overcome the methodological challenges posed by WhatsApps end-to-end encryption through a large-scale, privacy-preserving data donation infrastructure. Unlike public platforms (e.g., X/Twitter or Facebook) where data can be scraped or accessed via APIs, WhatsApps architecture necessitates a user-centric collection approach.

In this piece, we thus rely on data from a large WhatsApp data donation program that took place in Uttar Pradesh (India) in 2024, during the general elections campaign. While a fuller description of the tool we used to collect the data, WhatsApp Explorer - can be found in [Garimella and Chauchard \(2025\)](#), we note that the data collection underwent thorough ethics review and received approval at multiple institutions, as well as through the funder’s own ethics review process [Funders name, grant number and ethics review reference]. In summary, we recruited a cohort of 3,547 participants across 9 districts of the state of Uttar Pradesh (UP hereafter), India. We focused on UP as it is Indias most populous state and a critical political battleground that often serves as a bellwether for national electoral trends. Participants were recruited through on-the-ground partnerships to ensure the sample represented a broad sociodemographic spectrum. The cohort was stratified to capture diversity across key demographic axes, including age, income levels, and caste groups. This stratification is vital for ensuring that our analysis of information consumption captures the experiences of the general electorate rather than being skewed toward tech-savvy or urban elites.

Recruitment and participant characteristics. Recruitment was conducted by trained field enumerators working with two local survey partners, who approached eligible adults (18+, WhatsApp users) at the district level following a stratified sampling frame over age, gender, caste group, and income quintile. Enumerators explained the study in Hindi, walked participants through the informed-consent process, and assisted with on-device installation of WhatsApp Explorer (see [Garimella and Chauchard \(2025\)](#) for additional details). Selective opt-out is a genuine concern: participants who perceived their groups as highly sensitive (e.g., closed partisan organizers, journalists’ backchannels) were free to decline or to exclude specific groups at the source.

Our sample therefore likely under-represents elite political operatives and over-represents ordinary, “everyday” participants in public-facing group life. We discuss the implications of this opt-out structure for external validity in the Limitations subsection.

Group-level metadata. Each donated group is associated with metadata including the number of members at the time of donation, self-reported group type (family, neighborhood/-community, religious, political/party, news/updates, professional, entertainment), and, where available, the participant’s relationship to the group’s creation (member vs. admin). Across the 6,500 unique groups, the modal group size is 1,024 members (median 81, 90th percentile 521). Group-type composition is mixed: the majority are classified by participants as community/neighborhood or religious groups, with explicitly political/partisan groups a substantial minority. Detailed group-size and group-type distributions are reported in Appendix G, along with the fraction of AI content observed in each bin.

Under our protocol, consenting participants could only donate data from the large WhatsApp groups—groups of 4 members or more—that they were members of (one-on-one conversations, because they were likely to be of a more private nature, were by default excluded from the donation). WhatsApp Explorer tool was designed with a “privacy-by-design” architecture (Garimella and Chauchard 2025). The tool also performs extensive anonymization locally on the participants’ devices before any data was exported to our secure servers. This on-device processing included the removal of personally identifiable information such as phone numbers, names, and email addresses. Furthermore, to protect the privacy of non-consenting group members, faces appearing in images were automatically detected and blurred at the source. This rigorous protocol ensured compliance with ethical standards while allowing us to analyze the content circulating in semi-public digital spaces.

Data collection spanned the critical window of the Indian general election, from March to July 2024. This time frame captures the pre-election campaigning intensity, the polling phases, and the immediate post-election period. The resulting corpus comprises approximately 6,500 unique WhatsApp groups, yielding a total of 5.5 million messages. Given that generative AI is most frequently discussed in the context of *visual* misinformation and synthetic media, our analysis specifically isolates the 2.1 million images contained within this dataset.

Another crucial feature of our dataset is the preservation of metadata regarding mes-

sage dissemination. WhatsApp utilizes specific tags to indicate how content travels across the platform. Messages are flagged as “Forwarded” if they have been passed along once, and “Forwarded many times” if they have been forwarded through a chain of five or more hops. This distinction allows us to differentiate between organic, low-level sharing and content that has achieved significant virality. By retaining these indicators, we can weight our analysis toward content that actually reached a wide audience, distinguishing between high-visibility political communication and niche or ephemeral content.

3.2 Detecting Generative AI Content

Given the fear with which some observers were approaching the 2024 elections, one might hypothesize a high volume of synthetic media circulating in our dataset. To empirically test this, we focused our detection efforts primarily on image content. This decision is grounded in recent findings that the vast majority (approximately 76%) of AI-generated content on social platforms is image-based rather than video-based (Corsi, Marino and Wong 2024). Furthermore, qualitative studies of political campaigning infrastructure suggest that political propagandists continue to prioritize static imagery, such as meme templates, posters, and greeting cards, over complex video generation, due to ease of production and bandwidth constraints in the Global South (Daepf et al. 2025).

While recent work in the U.S. context has demonstrated that state-of-the-art Vision Language Models like GPT-4o can identify generative AI images with high accuracy using simple zero-shot prompting (Chen et al. 2025), applying this method directly to our entire corpus of over 2 million images was computationally and financially prohibitive. Therefore, we designed a multi-step, cascading detection pipeline designed to maximize recall in the early stages and precision in the later stages.

We began by selecting a random 10% sample of the total image dataset, resulting in approximately 210,000 images. To process this volume efficiently, we employed the open-weights Gemma 3 27B model as a first-pass filter. This model was tasked with flagging any potential AI-generated content. This initial sweep identified 10,500 images as potentially AI-generated. These flagged images were then subjected to a more rigorous verification step using the GPT-4o

model, which currently represents the industry standard for multimodal understanding. This second stage filtered the flagged set down to 3,158 images. Finally, to eliminate subtle false positives that automated systems often miss, such as hyper-edited “Good Morning” messages or traditional digital art common in Indian WhatsApp groups, three independent expert annotators manually reviewed the remaining 3,158 images. Aggregating their judgments by the median of their 1–5 AI-likelihood ratings, approximately 43% of these images (roughly 1,370) were classified as probable-AI (median ≥ 4), and approximately 18% (570 images) as high-confidence AI (all three annotators rated ≥ 4). Unless otherwise noted, all downstream analyses use the probable-AI subset, with robustness checks at the high-confidence threshold reported in the Appendix.

For the automated stages of the pipeline, we utilized a zero-shot prompt designed to distinguish between synthetic generation and standard digital manipulation. Crucially, the prompt included instructions to account for the privacy-preserving modifications (face blurring) introduced during our data collection process, preventing the models from misclassifying anonymization artifacts as AI “glitches.” The prompt used was:

Analyze this image and determine if it was created entirely by an AI image generation model (like DALL-E, Midjourney, or Stable Diffusion). Answer with a single word first: ‘yes’ if the image appears to be fully AI-generated, or ‘no’ if it appears to be a real photograph, screenshot, text message image, clipart, poster, meme, or other non-AI-generated content. Sometimes the faces are blurred due to anonymization so don’t be fooled by that.

This rigorous three-stage process moving from open-weights filtering to proprietary model verification to human validation yielded a directly detected prevalence rate of approximately 0.7% within our sample, which rises to about 1.1% once corrected for the measured recall of the first-pass filter (see below). In and of itself, this figure stands in stark contrast to the enchanted deterministic narratives that predicted an overwhelming flood of AI-generated content.

Concrete detection examples. To illustrate the behavior of the pipeline, we provide a small set of examples in Appendix E. In the positive direction, typical “yes” decisions involve images whose surfaces show the characteristic hyper-smooth, hyper-saturated finish of diffusion models (e.g., devotional composites of deities and lotus backgrounds), accompanied by malformed text glyphs, incoherent jewelry or weaponry, and uncanny hand anatomy. Typical “no” decisions include (a) human-photographed poster copies with JPEG compression and

photographic grain, (b) screenshot collages built in WhatsApp-native editing tools, and (c) face-blurred photographs from our privacy-preserving collection pipeline, where the prompt was explicitly engineered to prevent blurring from being misread as an “AI glitch.” Ambiguous cases—typically traditional non-AI digital art in the Indian devotional style—were the most common source of GPT-4o false positives and are precisely the cases human review was designed to catch.

Detection bias and the “absent-but-undetectable” concern. A reasonable criticism is that our pipeline detects AI primarily via visible artifacts and therefore risks systematically missing photorealistic, near-indistinguishable synthetic images, that is, the very deepfakes most feared by pre-election commentators. We take this concern seriously, and address it first where it bites hardest. Stage 1 is the decisive point: because it is a recall-oriented filter whose discards never reach the later stages, a miss there cannot be recovered downstream. We therefore validated the Gemma 3 filter directly, in two ways (Appendix D). We (i) benchmarked the exact Stage-1 prompt against a labeled AI-vs-real reference set, and (ii) audited a random sample of the images Gemma discarded, using an independent classifier (GPT-4o) followed by manual review of every flagged image. On the external benchmark, Stage-1 recall varies widely across generators, from 0.98 (Stable Diffusion XL) and 0.96 (DALL-E 3) down to roughly one-half for the most photorealistic recent models (Midjourney v6, 0.53; FLUX.1, 0.49) and for heterogeneous community art (0.56), pooling to 0.70 at uniformly high precision (≥ 0.95); recall is weakest exactly where synthesis artifacts are subtlest. The in-corpus audit measures that margin on our own data: five of 1,000 discarded images (0.5%) were confirmed AI on manual review, implying a corpus Stage-1 recall near 58% and a recall-corrected prevalence of about 1.1%, up from the 0.7% detected directly. Because this rests on five observed misses, it carries a wide interval: the 95% Wilson interval on the miss rate corresponds to a corrected prevalence of roughly 0.9–1.8%. We therefore report 1.1% as our central estimate and treat it as a lower bound, with the upper limit of that interval as the cautious reading.

The later stages are less exposed to this bias, though not immune to it. The intermediate stage (GPT-4o) is a frontier multimodal model benchmarked in prior work to recover AI images whose artifacts are not visible to untrained human coders (Chen et al. 2025), and our per-item human review by three independent annotators catches borderline high-fidelity cases that GPT-

4o flagged only weakly. No stage can rule out an entirely invisible AI corpus, but first-pass recall is now bounded empirically rather than assumed.

On a focus on images rather than AI-generated text. Our pipeline targets images rather than text. Textual AI content (e.g., ghostwritten campaign messaging) is in principle identifiable through stylistic signals (overuse of em-dashes, characteristic phrasing, register shifts), but reliable text-provenance classifiers remain an active research area and none has been validated against Hindi/Bhojpuri/code-mixed campaign text at WhatsApp scale. We therefore follow prior work in restricting our prevalence claims to images (Corsi, Marino and Wong 2024; Chen et al. 2025), where model confidence is markedly higher, while acknowledging that AI-authored text is a plausible but presently unmeasured channel of influence.

On the use of 2024-era models. The generative models available during our collection window (DALL-E 3, Midjourney v6, Stable Diffusion XL, and contemporaneous variants) had specific failure modes, including glossy surfaces, malformed text, and anatomical errors, that our human annotators were trained to detect. A natural question is whether rerunning the detection stages against newer frontier models (e.g., GPT-5 or successor Claude releases) would change our estimates. We expect that it would change them only modestly, and we think it is worth explaining why before describing the validation work we have actually undertaken. A stronger Stage-2 model is most likely to disagree with GPT-4o on borderline cases, that is, on images that GPT-4o was already uncertain about, and to push a small number of additional images into our human-review pool. Because every image that survives the automated stages is then independently judged by three trained annotators, those additional borderline images would simply join the human-review pool rather than directly change a final label. Their plausible net effect on our overall prevalence and content-type estimates is small (we would expect at most a one- to two-percentage-point increase in the share of images classified as “probable AI,” and no qualitative change in the dominant phatic, devotional, and greeting patterns we document). Newer frontier models are also released every few months, so any single such comparison risks being outdated by the time the paper appears in print. The more consequential question, whether the first-pass filter misses AI content outright, is addressed by the direct Stage-1 validation described above.¹

¹A dedicated methodological benchmark of frontier-model detection on Indian vernacular visual content would be a valuable companion study, and we view it as a natural follow-up project.

3.3 Annotation Framework and Protocols

Following the automated detection pipeline, the final set of flagged images underwent a granular manual annotation process. We developed a comprehensive coding instrument designed not only to verify the presence of generative AI but to understand its qualitative characteristics, deceptive potential, and functional utility in political discourse or in other types of communication. A specialized annotation portal was built to enable contextual review: annotators could examine each of the 3,100 pieces of content alongside the 10 messages sent immediately before and after it, preserving the conversational context in which the content appeared.

Detecting modern generative AI content and characterizing its communicative function requires a high degree of specialized visual literacy: the artifacts of AI, such as inconsistent lighting textures, subtle anatomical errors, and characteristic “glossy” surfaces, are often imperceptible to untrained observers. Given this inherent uncertainty, we rely on a three-annotator protocol and report inter-annotator reliability on every field used downstream.

Annotator panel. The three annotators were recruited from within a specialized research team that works on political communication and digital forensics in the Indian context. All three are fluent in Hindi and English, familiar with the visual vernacular of Indian WhatsApp (devotional art, political propaganda, “Good Morning” greeting conventions), and received a structured training module that included: (i) a walk-through of the annotation portal; (ii) annotated exemplars at each level of the 1–5 likelihood and quality scales; and (iii) a calibration batch of 100 images coded collectively before independent annotation began. After this training, each of the 3,158 flagged images was then independently annotated by all three coders. The full codebook, training exemplars, and screenshots of the annotation interface are provided in Appendix B.

Aggregation rules. For the per-image record used in all downstream analyses, we aggregate annotator judgments as follows: (i) Likert-scale fields (AI likelihood, quality, counterfactual necessity) are aggregated by median; (ii) binary classifications (e.g., *political*, *partisan*, *religious*, *hateful*, *misinformation*) by majority vote (2 of 3); (iii) multi-select categorical fields (e.g., content types, phatic sub-tags, partisan target) by majority endorsement of each tag; and (iv) free-text justifications are concatenated with annotator identifiers preserved. We additionally construct

a five-level consensus label (*high-confidence AI, probable AI, ambiguous, likely non-AI, uncertain*) based on the pattern of three likelihood ratings, and report robustness analyses under alternative thresholds in Appendix C.

Inter-annotator reliability. We report percent agreement, Fleiss’ κ (multi-rater generalization of Cohen’s κ) for binary fields, and Krippendorff’s α with an ordinal distance metric for Likert fields. Full IRR tables are in Appendix C. Briefly: for binary categorical judgments, agreement was substantial to near-perfect on tags tied to visible content ($\kappa = 0.81$ for *religious*, 0.78 for *socially relevant*, 0.81 for *advertising*, 0.60 for *phatic*, 0.60 for *partisan signal*), and moderate on judgments requiring political inference ($\kappa = 0.33$ for *political*). For Likert fields, exact-match agreement is lower by construction, but annotators fell within ± 1 scale point on 42% of *AI-likelihood* ratings and the median Likert aggregation is robust to an adversarial choice of any two annotators. These reliability patterns are consistent with the broader literature on expert content analysis of politically ambiguous visual material and support our decision to privilege the majority-endorsed aggregate over any single coder.

The annotation process evaluated each image across four distinct dimensions. The full codebook and questionnaire are detailed in Appendix B.

1. **Generative AI Confidence (1-5 Likert Scale):** The primary metric assessed the certainty of AI provenance. The scale ranged from 1 (“Definitely not AI-generated”) to 5 (“Definitely AI-generated”), with 3 indicating ambiguous cases where provenance was indeterminate. This scale allowed us to isolate high-confidence instances for our final prevalence statistics while retaining ambiguous cases for additional review.
2. **Production Necessity (Counterfactual Analysis):** We assessed the functional utility of the AI tool by asking: “*Could this image have been easily produced without AI?*” This dimension helps distinguish between revolutionary uses of the technology (e.g., creating photorealistic scenes of events that never happened) and efficiency-driven uses (e.g., generating generic images of deities). For instance, while AI was frequently used to generate idols, similar non-AI digital art is already ubiquitous in this ecosystem, suggesting AI served merely as a cost-saving measure rather than a productivity booster.
3. **Thematic and Rhetorical Analysis:** Beyond identifying the medium (AI vs. Human), we

systematically categorized the message content to understand its apparent communicative function. The annotator was tasked with summarizing the central claim or “main idea” of each image. This involved parsing the rhetorical structure of the message, specifically, whether the content appeared to be designed to inform, provoke, or solicit specific behaviors (e.g., voting, sharing, or purchasing). Because intent cannot be directly recovered without access to prompts or to the producer, we follow a cautious attributional frame throughout, inferring *supposed* or *apparent* communicative function rather than authorial intent (Brennen, Simon and Nielsen 2021). We also tracked observable social signals by coding for group reactions, allowing us to determine if specific types of AI-generated content elicited higher engagement or debate within the groups compared to benign or non-AI content.

4. **Epistemic Integrity and Harm Assessment:** To evaluate the potential for AI-driven disinformation, we employed a nuanced coding framework that transcends binary truth claims. We first assessed the *theoretical verifiability* of each image to distinguish factual assertions from subjective political satire. Beyond explicit falsehoods, we coded for *misleading framing* content that exploits omission or manipulation to deceive without being technically false and required qualitative justifications for such classifications. Finally, we screened all content for hate speech and toxicity to determine if generative AI was being weaponized against specific communities or remained confined to benign use cases.

This multi-dimensional framework allows us to move beyond simple prevalence counts to construct a taxonomy of *how* and *why* AI is actually being deployed in the wild.

4 Results

Our analysis of the dataset challenges the prevailing narrative of an AI-dominated election. By rigorously filtering the 2.1 million images through our multi-step pipeline, we found that generative AI content constituted approximately 1.1% of the total image dataset. This is a recall-corrected figure: the pipeline detects 0.7% directly, and correcting for the measured recall of the open-weights first-pass filter raises the estimate to roughly 1.1% (95% Wilson interval 0.9–1.8%;

Appendix D). We treat 1.1% as a lower bound, since the images least likely to be caught at any stage are the most photorealistic ones. This low prevalence serves as the foundational context for the detailed breakdown of content types, quality, and virality presented below.

4.1 Prevalence and Quality

Despite the availability of sophisticated tools, the volume of AI-generated content remained negligible. As shown in Figure 1, detecting this content was not a binary task. A significant portion of images fell into a “grey zone” (ratings 2-3), exhibiting some aesthetic markers of AI, such as hyper-saturated colors or the hyper-smooth, plastic-like surface finish we refer to as “glossy” throughout (see Appendix E for labeled exemplars), but lacking definitive artifacts; borderline and higher-fidelity cases such as these are what motivate the direct validation of first-pass recall reported in Appendix D. Using the median of our three annotators’ 1–5 AI-likelihood ratings, approximately 43% of the flagged set received a “probable AI” rating (≥ 4), and approximately 18% a “high-confidence” rating (all three annotators ≥ 4). Hereafter, “AI posts” refers to the probable-AI subset (median likelihood ≥ 4).²

Figure A4 (Appendix) presents representative examples across the five-point likelihood scale. A qualitative inspection reveals a striking visual convergence between traditional digital media and AI outputs. Across all rating levels, the images share a consistent vernacular aesthetic (or ‘vibe’) characterized by religious iconography, floral motifs, and greeting text overlaid on saturated backgrounds. Consequently, high-confidence AI images (Ratings 4-5) are often distinguished not by their subject matter, but by subtle technical artifacts such as incoherent text rendering, anatomical anomalies, or a distinctive hyper-smooth surface texture that deviates from standard digital editing.

Furthermore, among the content identified as definitely AI-generated, the technical sophistication was notably low. Figure 2 illustrates the quality ratings of confirmed AI images. The majority of images were rated as “poor quality” (1-2 on the quality scale), characterized by obvious artifacts, anatomical errors, and cartoonish aesthetics. This contradicts the “deepfake” narrative which assumes AI content is indistinguishable from reality; in the Indian context, AI content was often conspicuously artificial, though this observed skew toward low quality is

²Analyses at the high-confidence threshold appear in Appendix F.




Figure 1: Distribution of manual annotation ratings for AI likelihood (1=Definitely Human, 5=Definitely AI).




Figure 2: Quality ratings of confirmed AI-generated content (Likelihood ≥ 4).

partly a property of artifact-based detection, which more readily flags conspicuously synthetic images than the higher-fidelity output it misses (Appendix D). A caveat on the “quality” label is however in order: a cartoonish or stylized aesthetic is not intrinsically low-quality in a cultural setting where devotional illustration, political posters, and greeting-card design have long relied on bold, idealized, non-photorealistic styles. Our quality rubric measures technical execution given the producer’s evident stylistic goal (e.g., anatomical coherence, text legibility, lighting consistency) rather than aesthetic preference. A highly cartoonish image can still receive a high quality rating if the producer’s apparent goal was cartoon, and a photorealistic target with visible artifacts receives a low rating.

In sum, these findings challenge the narrative of an AI-saturated information ecosystem. Not only was the overall prevalence of generative AI negligible, but its application was also overwhelmingly banal. Rather than enabling novel forms of manipulation, the technology was primarily co-opted to reproduce existing visual tropes, specifically religious iconography and phatic greetings, serving essentially as a high-tech substitute for traditional digital collages. In

the next section, we dig deeper into specific use cases of generative AI content.

4.2 A Taxonomy of AI Content During The 2024 Elections Campaign

If AI was not being used to create high-fidelity deepfakes, what was it used for? To answer this, we developed a functional taxonomy based on the primary communicative intent of each image (see Appendix for the full coding instrument). We classified content into distinct categories ranging from social maintenance (the phatic register we define below) to political persuasion.

Figure 3 presents the distribution of these categories. The ecosystem is overwhelmingly dominated by what we term “phatic” content (67%). Drawing from anthropological definitions of phatic labor, we define this category as content designed primarily to build or maintain social relationships rather than to convey specific information this includes the ubiquitous “Good Morning” greetings, festival wishes, and generic devotional imagery.

The second largest category was “Art/Entertainment” (23%), comprising aesthetic digital art, memes, and humorous content devoid of explicit political messaging. In contrast, instrumental categories such as “Logistical” (ads, practical information) or malicious categories such as “Scams” and aggressive political propaganda constituted a minor fraction of the dataset. This distribution confirms that for the average user, generative AI served largely as a tool for social signaling and aesthetic expression rather than information warfare.

A deeper analysis of the phatic category, shown in Figure 4, reveals that approximately one-third of all AI images circulating in our dataset were generic “Good Morning” messages or greetings. Note that the three bars in Figure 4 do not sum to the full ~66% phatic share reported above, nor to 100%: they show the share of *all AI posts* that carry each specific sub-tag, and most of the remaining phatic content consists of generic devotional or festival imagery that was tagged simply as “phatic” without a more specific sub-label. The three sub-tags shown are therefore a partial, not exhaustive, decomposition of the phatic category. This suggests that the primary driver of AI adoption in this election cycle was not deception, but *efficiency*. Users and content creators utilized AI tools to generate quick, visually appealing greeting cards and inspirational quotes, as they presumably had done before the electoral campaign (Chauchard and Garimella 2022), rather than to manipulate public opinion. While we detect very little content that can




Figure 3: Distribution of AI content by category.

credibly be classified as fulfilling this goal, we note that this number would in all likelihood have been even smaller before and after the end of the campaign, as interest in politics waned.

Table A1 (Appendix) further corroborates this, showing that high-confidence AI images (Score 5) were disproportionately concentrated in the “Phatic” (73.3%) and “Art” (31.8%) categories, whereas lower-confidence images were often logistical or advertising spam. Visual examples of this benign content such as AI-generated flowers or generic greetings can be seen in Figure A6 in the Appendix.

4.3 Political and Social Signals

While the overall volume of political AI content was low, its distribution reveals significant asymmetries. To distinguish between general cultural content and explicit electioneering, our coding scheme (detailed in the Appendix) differentiated between three categories of relevance. We defined “Socially Relevant” content as images pertaining to general associational life, such as festivals, cultural symbols, or deities, without necessarily invoking political actors. In contrast, the “Political” tag was reserved for content discussing governance or public affairs, while “Partisan”




Figure 4: **Breakdown of phatic AI content**, dominated by greetings and motivational quotes.

was strictly applied to messages explicitly promoting or attacking a specific party or leader. Our “partisan” category is thus nested within our “political” category: every *partisan* appeal is by construction *political*, but not every political image names a party or leader (e.g., generic appeals to vote, abstract references to governance, or institutional imagery). Coding them as separate, non-exclusive fields allows us to distinguish the general politicization of content from explicit party-level mobilization.

As shown in Figure 5, the results are stark. Over half (54%) of the AI content was broadly “Socially Relevant,” reflecting the technology’s use in cultural signaling. Conversely, only 3.7% was explicitly “Political,” and a mere 3.1% met the strict criteria for “Partisan” propaganda.

A caution on the “Social” / “Political” boundary. We read the “social $\gg$ political” pattern with caution. In the Indian setting—and especially during the 2024 campaign period—devotional and patriotic imagery are not apolitical: Hindutva-coded deities, cow imagery, “love jihad” framings, and hyper-saturated renditions of national monuments carry well-documented political resonance (Mehta 2024; Shih 2023). Our strict coding scheme deliberately requires explicit reference to a party, leader, election, or policy issue before a *Political* or *Partisan* tag is assigned, which is conservative by design. We therefore do not claim that socially-coded AI con-

tent is politically inert; we claim only that its explicit political or partisan content is low, while recognizing that identity-activating imagery can shape political attitudes through mechanisms that our coding scheme is not instrumented to detect. Measuring whether and how such content translates into attitudinal change would require an accompanying panel survey of recipients, an important direction for future work that we flag explicitly in the Limitations subsection. Our prior, however, is that the type of phatic content we have documented above should *not* lead to either sudden or large shifts in political preferences or majorities.




Figure 5: Proportion of AI content containing social, political, or partisan signals.

We also examined the “habitats” of this content to understand the contexts in which it circulated. Figure 6 reveals that nearly 50% of AI images appeared in Religious groups, followed closely by “Political” (31%) and “Hindutva/BJP” (25%) groups. Read against the corpus-wide composition of group types (Appendix G), where religious and Hindutva/BJP groups make up 19.7% and 9.1% of all groups respectively, AI images are roughly twice as concentrated in these spaces as their baseline share would predict. This distribution highlights a significant functional dynamic: while the AI content itself is rarely explicitly political (as shown in Figure 5), it is heavily utilized within political and partisan infrastructures. This suggests that benign AI-generated imagery, such as religious idols and morning greetings, plays a critical role in the

maintenance of ties and in sociality between group members. Prior research had shown that in these political spaces, high-aesthetic, low-stakes content likely serves to keep groups active and socially lubricated, maintaining user engagement and communal bonding during lull periods between explicit political mobilization (Chauchard and Garimella 2022). We show here that this extends to campaign periods during which more openly political, or even partisan, content would be expected to dominate, as well as to the specifically AI-built subset of the content. Despite the proximity of an election happening in a context of extreme polarization in India, users of these threads had recourse to AI technology first and foremost to salute one another and send wishes.




Figure 6: Distribution of AI content across different WhatsApp group types.

4.4 Virality of AI content


A recurring concern is that even small absolute volumes of AI-generated imagery could have outsized influence if they go viral. WhatsApp preserves a “forwarded many times” flag for messages that have traveled through a chain of five or more hops, which allows us to directly compare the virality of AI vs. non-AI images in the same groups. Figure A1 (Appendix) shows that AI images are in fact *less* likely than non-AI images to reach the “forwarded many times”

state, and their peak forwarding counts are similarly lower. This is consistent with our demand-side interpretation: the glossy synthetic aesthetic plausibly functions as a low-credibility heuristic among recipients, dampening onward sharing. Even so, we note three qualifications. First, “low virality on average” does not rule out a small number of high-impact AI images; our corpus does contain individual high-forward AI cases, which we catalogue in the Appendix. Second, this pattern is specific to 2024-era detectable-AI aesthetics; as future models approach photorealism, the virality penalty we document may dissipate. Third, WhatsApp’s forwarding-hop counter is a coarse proxy for cross-group reach, since highly duplicated content produced in parallel by many accounts does not accumulate hops in the same way.

4.5 Misinformation and Hate

This relative absence of political content within the AI-build subset of our data drives another surprising result. Contrary to predictions of deepfake-driven chaos, misinformation and hate speech were indeed almost non-existent in our AI subset. Hate speech constituted approximately 1% of the AI images.


A qualitative analysis of the rare misinformation examples (Figure 7) reveals that they seldom constituted malicious political fabrications in the conventional sense. Instead, they were primarily instances of “myth-making,” characterized by pseudoscientific health claims (e.g., the “Sanjeevani Mantra” (translation: Liferestoring chant) for longevity), exaggerated religious miracles, or hyper-idealized, futuristic renderings of public infrastructure. Similarly, the sparse instances of hate speech (Figure 8) were technically crude and relied heavily on established sectarian tropes, specifically narratives regarding “Love Jihad” or “Population Jihad”, rather than novel, AI-generated falsehoods. Consequently, even in those rare cases when it was arguably employed to shape political or communal beliefs, the technology was not deployed to engineer *new* forms of deception, but rather to visually amplify existing folk theories and social prejudices.



(a)




(b)



कंबोडिया समुन्द्र के बीचो-बीच जैन
 धर्म के तेइसवें तिर्थंकर श्री पार्श्वनाथ
 प्रभु की प्रतिमाजी पहाड को त
 हजारो साल पूर्व उस समय के
 राजा द्वारा बनवायी गयी थी।

(c)


Figure 7: Misinformation examples



(a)



(b)



(c)

Figure 8: Hateful examples

5 Discussion

Our investigation into the 2024 Indian General Election offers a rejoinder to the “enchanted determinism” that characterized pre-election discourse (Campolo and Crawford 2020). The dominant narrative, propelled by security agencies, media outlets, and technology vendors, posited that generative AI would act as an unstoppable, exogenous shock to the democratic system, flooding the zone with undetectable deepfakes and personalized propaganda.

The empirical reality was different. With generative AI constituting merely 1.1% of the image dataset, the anticipated “tsunami” never materialized. Even when generative AI was used, it was just to complement existing workflows, e.g. to create memes and phatic content. This disconnect forces us to re-evaluate the incentives and constraints governing political communication in the Global South. The failure of AI to transform the 2024 election was not merely a matter of technological immaturity. The practices of political parties (such as the ruling BJP) in terms of content creation are likely inelastic in the short term — costs and skills acquisition constitute significant hurdles. Besides, there is no reason to expect that the apparition of a new technology would change the preferences of users in terms of content. Insofar as phatic, rather than political or partisan content, has dominated and continues to dominate social media posting in India, it would have been surprising to see *political* deepfakes take over. Logically enough, AI was used, however disproportionally so by creators of phatic content.

5.1 The Friction of Reality: Why Predictions Failed

The gap between the predicted “AI Election” and the ground truth can be attributed to what we term *technological friction*. The enchanted view of AI assumes that because a technology *can* be used for deception, it *will* be. However, this ignores the logistical realities of political actors that would most likely be in charge of producing such content in India, the famous “IT cells” of parties, for instance (Mehta 2024). For a political operative, the goal is not to use the most advanced technology, but to use the most effective one. Creating a high-fidelity deepfake that survives scrutiny requires significant prompt engineering, consistency checks, and often post-processing is a high-friction workflow. In contrast, “cheapfakes”—miscontextualized videos, edited newspaper clippings, or traditional text-based rumors—are low-friction, cheap to produce,

and historically proven to drive engagement (Dan et al. 2021). Our findings suggest that AI had, as of 2024 and in India, not yet crossed the threshold of utility after which it is more efficient than a cheapfake for the specific purpose of negative campaigning. The “dog that didn’t bark” in this election suggests that disinformation actors are rational economic agents who prioritize return on investment over technological novelty. For instance, see Figure A5 shows examples of misinformation images which were not generated using generative AI. We can clearly see most examples are simple images with text on a plain background or on an image, or image collages, or simple animations.

It is crucial, however, to contextualize these findings from the technological vantage point of 2024. The models available during our data collection window (e.g., Midjourney v6, DALL-E 3) still struggled with text rendering and specific Indian cultural markers. While the technology has and will undoubtedly improve, our data may suggest that the social demand for high-fidelity deception is lower than assumed; in a context of high social and partisan polarization, human-generated misinformation may suffice.

Finally, it is imperative to acknowledge that the disparity between prediction and reality does not necessarily invalidate the concerns raised by the expert community. Indeed, the high-volume discourse surrounding the “AI Election” may have functioned as a self-defeating prophecy in the most positive sense. The widespread media coverage and pre-bunking campaigns led by civil society organizations likely increased digital literacy, priming the electorate to be hyper-vigilant. By creating a climate of suspicion around synthetic media, these predictions may have inadvertently raised the bar for successful deception, rendering low-quality deepfakes ineffective before they could proliferate.

5.2 Supply-Side Dynamics: Automation, Not Deception

If AI was not used to deceive, what was it used for? Our finding that the majority of AI content falls into “Phatic” and “Religious” categories reveals that Indian users have adopted AI primarily as a *cost-saving mechanism* to keep producing pre-existing contents in a more efficient manner. In the high-volume environment of Indian WhatsApp, the exchange of “Good Morning” messages and religious greetings is a vital ritual for maintaining social ties. Generative AI, with its

tendency toward hyper-saturated colors, symmetrical composition, and idealized imagery, fits the existing aesthetic vernacular of this genre perfectly. AI was merely used to automate the production of the “glossy” aesthetic that Indian users already preferred.

Furthermore, our finding that 84% of partisan AI content (a small share of all AI content, as shown above) favored the ruling BJP challenges the “democratization” hypothesis (Kreps and Kriner 2023). Early optimists argued that AI would level the playing field, allowing resource-poor parties to generate high-quality content at near-zero cost. Instead, we observe that AI adoption mirrors existing offline power dynamics. The BJP’s dominance suggests that the bottleneck in digital campaigning is not *content production* (which AI solves), but *distribution networks* (which AI does not). The ability to push content to millions of voters still requires the massive human infrastructure and WhatsApp group networks that the ruling party has spent over a decade cultivating (Shih 2023). AI simply allowed the strong to become more efficient, rather than empowering the weak.

5.3 Demand-Side Dynamics: The Heuristics of Skepticism

The limited virality of AI content (Figure A1, Appendix) provides critical insight into audience reception. The fact that AI images had lower viral peaks than human content suggests that users may possess an intuitive skepticism (or perhaps indifference) toward the “synthetic” aesthetic. We hypothesize that the specific visual artifacts of 2024-era AI (plastic skin textures, incoherent background details) act as a heuristic signal for “low-effort” or “spam” content. Just as email users learned to ignore messages with poor grammar, WhatsApp users may be learning to filter out the glossy, uncanny look of AI images. Far from being the “perfect weapon” of influence, overt AI generation may currently impose a penalty on credibility. The content is recognized not as documentation of reality, but as “digital art” or “forwarded clutter,” reducing its potential to mobilize political anger or belief.

5.4 Limitations

Our study has several limitations that shape the appropriate interpretation of its findings.

Sampling and representativeness. Our data come from consenting participants in UP

who agreed to donate messages from groups of four or more members. While UP is electorally decisive, and the largest Indian state, it is not the entirety of India; cross-state variation in language, platform mix, and partisan infrastructure means that our prevalence estimates should not be read as national point estimates but as the most detailed available snapshot of a critical battleground state. Within UP, the opt-in nature of data donation plausibly under-represents participants who regard their WhatsApp activity as especially sensitive—campaign operatives, organized IT-cell volunteers, and members of closed partisan organizing groups—so our sample is closer to ordinary citizens’ information environments than to the production side of campaign messaging. By excluding one-on-one chats and small groups, we also cannot observe AI content in private couples, family, or close-friend messaging, where more personal (including romantic or explicitly partisan) content may circulate. Finally, broadcast-style channels (WhatsApp Channels, Telegram broadcasts) that campaigns are increasingly using for AI-augmented outreach fall outside our frame.

Detection pipeline. Our pipeline identifies AI primarily through visible synthesis artifacts, so the most photorealistic images are also the ones most likely to escape it; we bound this directly by validating first-pass recall and reporting a recall-corrected prevalence (Section 3.2 and Appendix D), but no correction of this kind can be exact. The residual bias is not politically neutral: an image that passes for a genuine photograph may well be *more* consequential than one recognizable as synthetic, so the material we are most likely to have missed is also the material most likely to persuade, and our estimate should be read with that asymmetry in mind. Closing this gap will require provenance-based approaches rather than better artifact detection alone, such as model-side watermarking and C2PA-style content credentials, combined with periodic re-benchmarking of detectors as generators improve.

Temporal scope. Our collection window (March–July 2024) reflects the capabilities of models available at that time. Generative image models have improved substantially since, in ways that specifically target the “glossy surface” and “malformed text” artifacts on which 2024-era detection relied. We therefore do not claim that 1.1% is a durable rate; the empirical question is whether improving supply has shifted *actual* WhatsApp circulation in 2025 and beyond, which would require new data collection.

Our design measures what is circulating, not what recipients do with it. We cannot

observe whether exposure to devotional AI imagery with implicit Hindutva framings changed political attitudes, affective polarization, or turnout. A properly powered test would combine our message-circulation data with either an experiment or with a pre/post panel survey of participants' attitudes—a design we flag as a priority for follow-up work.

Our strict “Political” / “Partisan” coding rules privilege explicit references over identity-activating imagery. We have already cautioned in Results that a “socially-coded” AI image with Hindutva symbolism is not politically neutral. An auxiliary coding layer that tags identity-salient imagery separately from party-explicit imagery would sharpen this distinction; we preview such a taxonomy in Appendix B as groundwork for future extensions.

5.5 Implications for policy and future work

These findings have context-bounded but meaningful implications for AI governance. Policy responses globally have focused on technical solutions like watermarking or detection to prevent deepfakes. However, our data suggests that in the 2024 Indian context, policymakers may be fighting a largely theoretical war while under-reacting to the very real, “old-fashioned” challenge of cheapfake circulation and long-running rumors (Paris and Donovan 2019; Brennen, Simon and Nielsen 2021; Badrinathan, Chauchard and Siddiqui 2025). Transparency and watermarking regimes may still be justified on grounds of principled public-interest disclosure even where prevalence is currently low, and their value will grow as detectable-AI aesthetics fade. But policymakers eager to protect vulnerable populations from manipulation and misinformation may want to allocate resources to tools and strategies to enhance defenses against more traditional formats already used by political actors. In India, rumors targeting minorities have for instance long been prevalent and recent analyses suggest they are currently widespread; So much so, that it is in fact unclear to us how even the most professionally produced deepfakes could aggravate this state of affairs.

Crucially, the relevance of these findings extends beyond the specific technical capabilities of the 2024-era models used in this study. Critics might argue that the limited impact we observed was merely a function of the “uncanny valley” artifacts present in tools like DALL-E 3 or Midjourney v6. However, this technological determinism ignores the enduring sociological and economic

constraints we identified. The friction against AI adoption was not just *fidelity*, but *utility*. Even as future models achieve perfect photorealism, the economic logic of the “cheapfake”—which creates high engagement with zero compute costs and minimal effort—remains a formidable competitor. Our study establishes a baseline demonstrating that without a shift in the *demand* for high-fidelity deception, improvements in the *supply* of AI capabilities may yield diminishing returns in actual electoral influence.

First, the obsession with “deepfakes” distracts from the ongoing prevalence of “shallow” misinformation and hate speech, which constituted the vast majority of problematic content in our dataset. Second, the real danger of AI in elections may not be the *quality* of the lie, but the *quantity* of the noise. By automating the creation of phatic and religious content, AI allows actors to flood information channels with benign noise, potentially drowning out genuine political discourse and exhausting user attention.

Finally, the enchanted deterministic narrative itself may pose a risk to democratic integrity. By continuously warning voters about “undetectable” AI, experts may inadvertently contribute to a “liar’s dividend,” where political actors can dismiss genuine incriminating evidence as “AI-generated” (Schiff, Schiff and Bueno 2025). Our study suggests that to preserve electoral integrity, we must pivot from fearing a hypothetical technology to understanding the mundane, efficient, and culturally specific ways it is actually being adopted.

We close with three forward-looking considerations. First, AI-generated *text*—ghostwritten slogans, templated WhatsApp messages, chatbot-mediated voter outreach—remains an unmeasured channel in our design, and is plausibly a larger actual vector of AI influence than imagery during the 2024 campaign. Reliable provenance detection for Hindi and code-mixed campaign text is an open research problem, and a follow-up wave of data donation paired with a text-provenance validation study is an immediate next step. Second, the prevalence we detect should not justify complacency: the distributional patterns we document (BJP dominance, concentration in religious and Hindutva-coded groups) imply that even a modest future increase in AI capacity would be absorbed by existing asymmetric infrastructures, amplifying rather than democratizing political communication. Third, because our design cannot directly test attitudinal effects, the strongest interpretation available from these data is about *prevalence and supply-side use*; claims about attitude formation, turnout, or “democratic resilience” require different designs and data.

We hope communication researchers will take up this challenge, as better understanding the attitudinal or behavioral effects of exposure to AI-generated content, however rare it might be, is necessary to gauge what resources should really be allocated to AI governance.

Data and code availability

Analysis code and the aggregated data needed to reproduce every figure, table, and reported statistic are deposited in a public repository ([\[INSERT REPOSITORY URL\]](#)). The deposit includes the annotation codebook and the per-image aggregated annotation records, the inter-annotator reliability computations, the group-type composition data, and the full Stage-1 validation materials: the benchmark script and its per-subset predictions and metrics, and the in-corpus false-negative audit code and results. Raw message content and the donated images themselves cannot be released. Participants consented to research use under a privacy-by-design protocol in which identifiers are stripped at collection, and the images frequently depict or reference third parties who never consented to disclosure; releasing them would breach both commitments. Researchers who need access to the restricted materials may contact the authors to arrange a data-use agreement.

References

- Badrinathan, Sumitra, Simon Chauchard and Niloufer Siddiqui. 2025. "Misinformation and support for vigilantism: An experiment in India and Pakistan." *American Political Science Review* 119(2):947–965.
- Bennett, W. Lance and Alexandra Segerberg. 2013. *The Logic of Connective Action: Digital Media and the Personalization of Contentious Politics*. Cambridge: Cambridge University Press.
- Bossetta, Michael. 2018. "The Digital Architectures of Social Media: Comparing Political Campaigning on Facebook, Twitter, Instagram, and Snapchat in the 2016 U.S. Election." *Journalism & Mass Communication Quarterly* 95(2):471–496.
- Brennen, J. Scott, Felix M. Simon and Rasmus Kleis Nielsen. 2021. "Beyond (Mis)Representation: Visuals in COVID-19 Misinformation." *The International Journal of Press/Politics* 26(1):277–299.
- Bueno de Mesquita, E., B. Canes-Wrone, A. B. Hall, K. Lum, G. J. Martin and Y. R. Velez. 2023. Preparing for Generative AI in the 2024 Election: Recommendations and Best Practices Based on Academic Research. Technical report Stanford Graduate School of Business & University of Chicago Harris School of Public Policy. Available at <https://www.gsb.stanford.edu/faculty-research/publications/preparing-generative-ai-2024-election-recommendations-best-practices>.
- Campolo, A. and K. Crawford. 2020. "Enchanted Determinism: Power without Responsibility in Artificial Intelligence." *Engaging Science, Technology, and Society* 6:1–19.
URL: <https://doi.org/10.17351/ests2020.277>
- Carr, Robert and Paul Köhler. 2024. Ai-pocalypse now? disinformation, AI, and the super election year. Munich Security Conference. <https://doi.org/10.47342/VPRS3682>.
- Chadwick, Andrew. 2017. *The Hybrid Media System: Politics and Power*. Second ed. Oxford University Press.
- Chauchard, Simon and Kiran Garimella. 2022. "What circulates on partisan WhatsApp in India? Insights from an unusual dataset." *Journal of Quantitative Description: Digital Media* 2.

- Chen, Zhiyi, Jinyi Ye, Beverly Tsai, Emilio Ferrara and Luca Luceri. 2025. Synthetic politics: Prevalence, spreaders, and emotional reception of AI-generated political images on X. In *Proceedings of the 36th ACM Conference on Hypertext and Social Media*. pp. 11–21.
- Christopher, N. and V. Bansal. 2024. “Indian voters are being bombarded with millions of deep-fakes. Political candidates approve.” *WIRED* . Available at <https://www.wired.com/story/indian-elections-ai-deepfakes/>.
- Corsi, Giulio, Bill Marino and Willow Wong. 2024. “The spread of synthetic media on X.” *Harvard Kennedy School Misinformation Review* 5(3):1–19.
- Cybersecurity and Infrastructure Security Agency. 2024. “Risk in Focus: Generative AI and the 2024 Election Cycle.” Online. Available at <https://www.cisa.gov/resources-tools/resources/risk-focus-generative-ai-and-2024-election-cycle>.
- Daepf, Madeleine IG, Alejandro Cuevas, Robert Osazuwa Ness, Vickie Yu-Ping Wang, Bharat Kumar Nayak, Dibyendu Mishra, Ti-Chung Cheng, Shaily Desai and Joyojeet Pal. 2025. “Generative Propaganda.” *arXiv preprint arXiv:2509.19147* .
- Dan, Virela, Britt Paris, Joan Donovan, Michael Hameleers, Jon Roozenbeek, Sander van der Linden and Christian von Sikorski. 2021. “Visual mis- and disinformation, social media, and democracy.” *Journalism & Mass Communication Quarterly* 98(3):641–664.
- Dobber, Tom, Nadia Metoui, Damian Trilling, Natali Helberger and Claes De Vreese. 2021. “Do (microtargeted) deepfakes have real effects on political attitudes?” *The International Journal of Press/Politics* 26(1):69–91.
- Elliott, V. 2024. “2024 is the Year of the Generative AI Election.” *WIRED* . Available at <https://www.wired.com/story/2024-is-the-year-of-generative-ai-elections/>.
- Escue, Brianna, Nathanael S Martin and Anthony Salerno. 2025. “That’s so cringeworthy! Understanding what cringe is and why we want to share it.” *Journal of Marketing Research* 62(4):664–683.

- Farooq, Aqsa, Elske van den Hoogen, Marina Tulin and Claes de Vreese. 2025. "Generative AI Generating Concerns: Citizens' Perspectives During the 2024 European Elections." *The International Journal of Press/Politics* .
- Ferrara, Emilio. 2024. "Charting the landscape of nefarious uses of generative artificial intelligence for online election interference." *arXiv preprint arXiv:2406.01862* .
- Freelon, Deen and Chris Wells. 2020. "Disinformation as Political Communication." *Political Communication* 37(2):145–156.
- Garimella, Kiran and Simon Chauchard. 2025. "WhatsApp explorer: A data donation tool to facilitate research on WhatsApp." *Mobile Media & Communication* 13(3):481–503.
- Gillespie, Tarleton. 2010. "The politics of 'platforms'." *New Media & Society* 12(3):347–364.
- Guess, Andrew M., Brendan Nyhan and Jason Reifler. 2020. "Exposure to untrustworthy websites in the 2016 US election." *Nature Human Behaviour* 4(5):472–480.
- Kreps, Sarah and Doug Kriner. 2023. "How AI threatens democracy." *Journal of Democracy* 34(4):122–131.
- Mehta, Nalin. 2024. How the Bjp Became the World's Largest Political Party: Organisational Restructuring and the Use of Digital Technologies. In *The New BJP*. Routledge pp. 168–217.
- Miller, Daniel, Jolynna Sinanan, Xinyuan Wang, Tom McDonald, Nell Haynes, Elisabetta Costa, Juliano Spyer, Shriram Venkatraman and Razvan Nicolescu. 2016. *How the world changed social media*. UCL press.
- Paris, Britt and Joan Donovan. 2019. "Deepfakes and cheap fakes." *Data & Society* 1.
- Rebelo, Karen. 2024. "Indias generative AI election pilot shows artificial intelligence in campaigns is here to stay." *Center for Media Engagement at The University of Texas at Austin* pp. 12–16.
- Schiff, Kaylyn Jackson, Daniel S Schiff and Natália S Bueno. 2025. "The liars dividend: can politicians claim misinformation to evade accountability?" *American Political Science Review* 119(1):71–90.

Shih, Gerry. 2023. "Inside Hindu Nationalists Vast Digital Campaign to Inflamm India." *The Washington Post* . News article.

URL: <https://www.washingtonpost.com/world/2023/09/26/hindu-nationalist-social-media-hate-campaign/>

Strikovic, Edina and Hannes Cools. 2025. "Reality Re-Imag (in) ed. Mapping Publics Perceptions and Evaluations of AI-Generated Images in News Contexts." *Digital Journalism* pp. 1–22.

Vaccari, Cristian and Andrew Chadwick. 2020. "Deepfakes and disinformation: Exploring the impact of synthetic political video on deception, uncertainty, and trust in news." *Social media+ society* 6(1):2056305120903408.

Valente, Thomas W. 1996. "Social network thresholds in the diffusion of innovations." *Social Networks* 18(1):69–89.

Votta, Fabio and Simon Kruschinski. 2025. "Campaign Tracker: Tracking Political Advertising and AI-Generated Content in European Elections." <https://www.campaigntracker.de/>. Accessed 2026.